


PARALLEL COMPUTING PARADIGMS IN THE CLOUD



PARADIGMS OF PARALLEL COMPUTING IN THE CLOUD

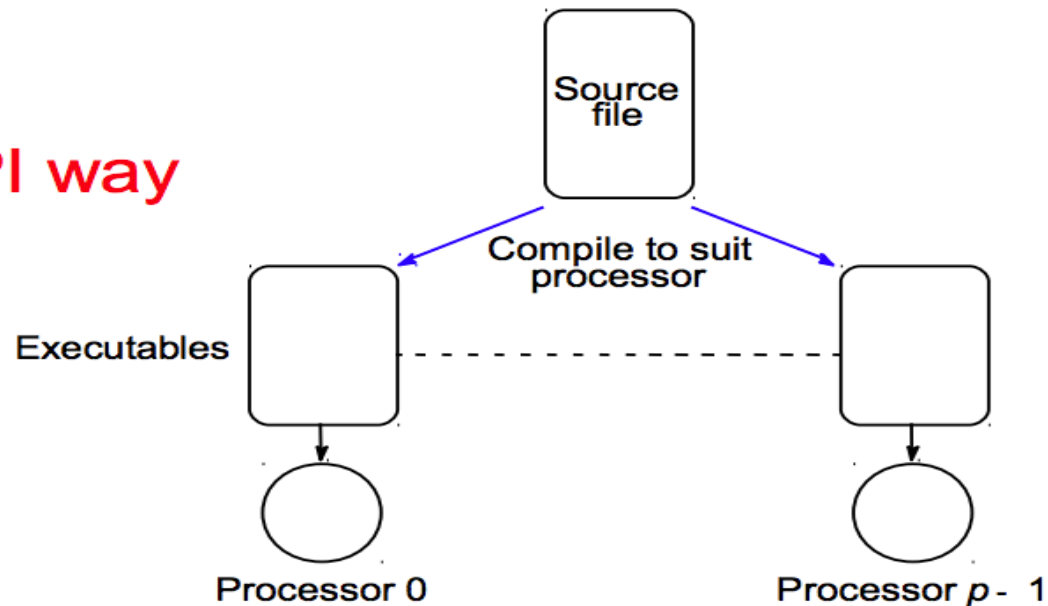
- Single Program Multiple Data (SPMD)
 - Classic High Performance Computing (HPC)
- Many Task Parallelism
 - A large queue or queues of tasks may be executed in any order
- Bulk Synchronous Parallelism (BSP)
 - Map Reduce
- Graph Execution
 - Spark and streaming systems
- Microservices
 - Computing is performed by one or more actors which communicate via message
- Serverless
 - Focus on application, not the infrastructure



SPMD

- Message Passing Interface (MPI) Programming Model
 - A computation comprises one or more processes that communicate by calling library routines to send and receive messages to other processes
- Same program executed by each processor
- Control statements select different parts for each processor to execute

Basic MPI way



A SIMPLE SPMD

```
#include <stdio.h>
#include <mpi.h>
#include <stdlib.h>

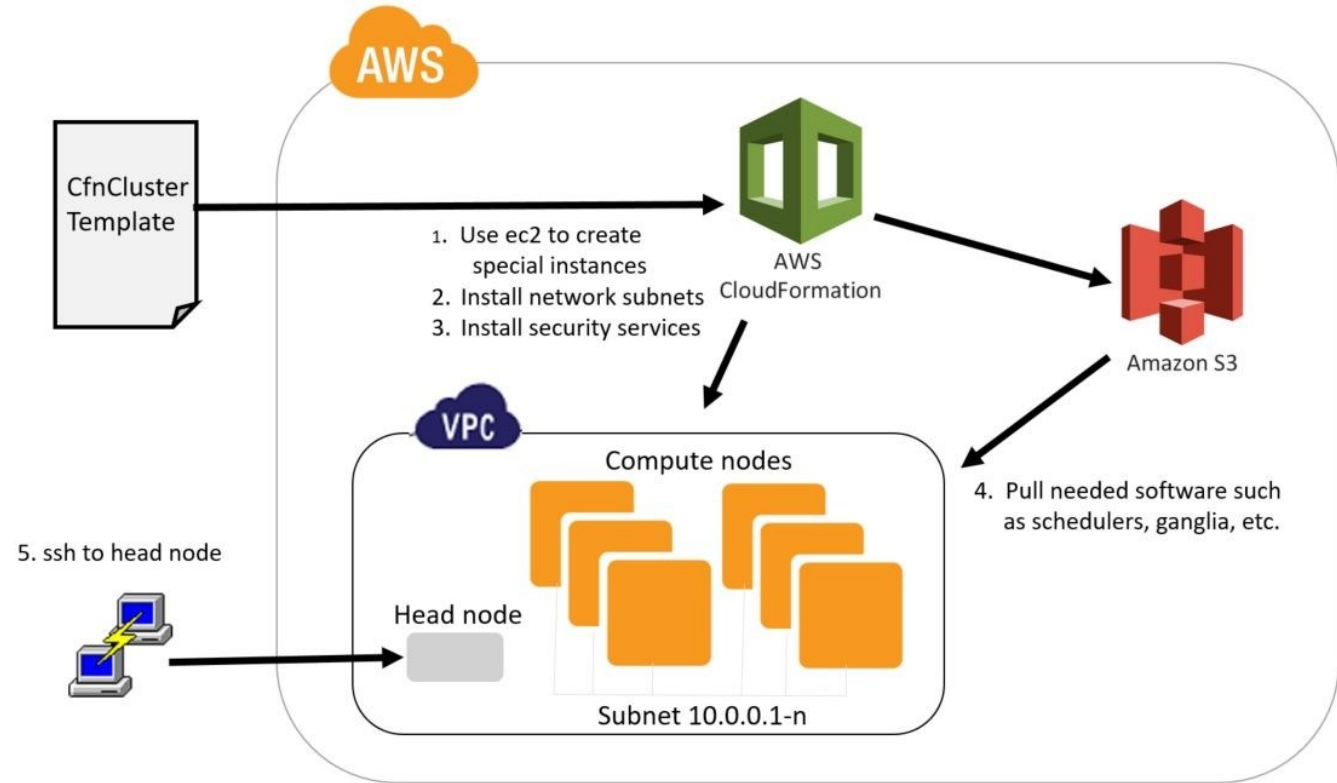
main(int argc, char **argv)
{
    char hostname[1024];
    gethostname(hostname, 1024);
    printf("%s\n", hostname);
}
```

```
mpicc ip-print.c
srn -n 16 /home/ec2-user/a.out > machines
```



SPMD USING AWS

- Amazon CloudFormation service
 - Enables automated deployment of complex collections of related services, ie.
 - Multiple EC2 instances
 - Load-balancers
 - Special network connecting them
 - Security groups
- AWS CloudFormation Cluster
 - Fill out CfnCluster template
 - Use aws command line to submit
 - Log into head node



SPMD USING AZURE

- Azure's Slurm Cluster service
 - Fill out the template similar to AWS
 - Setup a Slurm Cluster

Deploy a slurm cluster

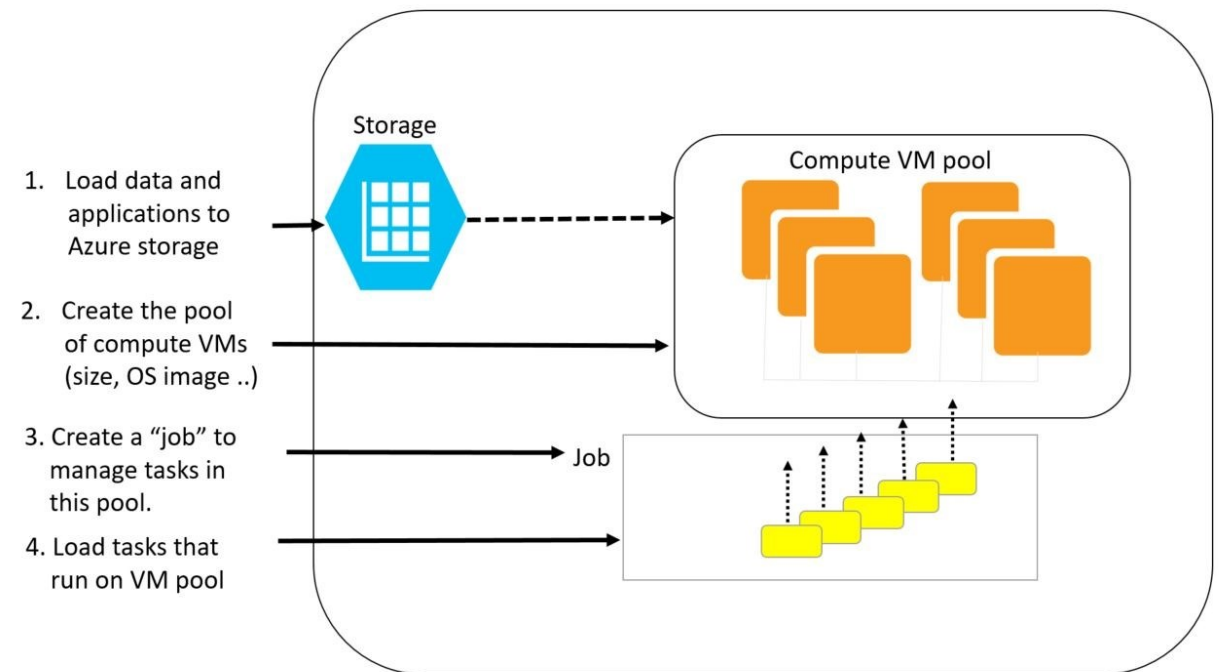


1. Fill in the 3 mandatory parameters - public DNS name, a storage account to hold VM image, and admin user password.
2. Fill in other info and click "OK".

Using the cluster


Simply SSH to the master node and do a srun! The DNS name is *dnsName.location.cloudapp.azure.com*, for example, *yidingslurm.westus.cloudapp.azure.com*.

- Use Azure Batch
 - Similar to AWS batch



MANY TASK PARALLELISM

- Task parallel model is great for solving problems that involve doing many independent computations




Each worker repeatedly pulls a sample from a queue, processes the data, and stores the result in a shared table



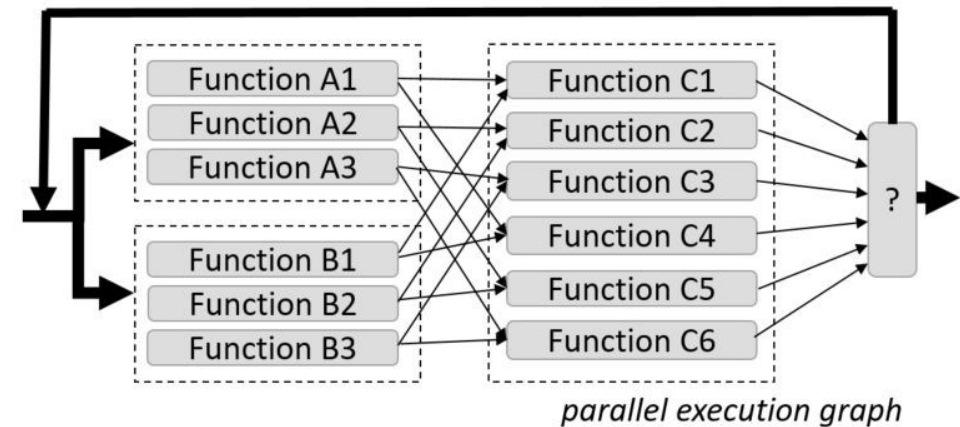
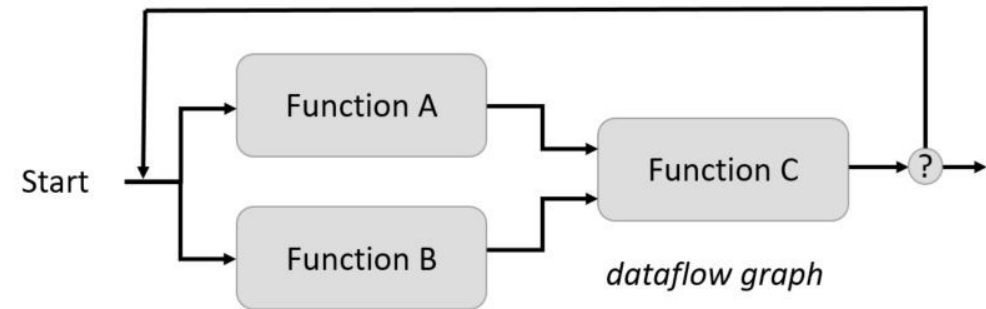
BSP — MAPREDUCE

- Worker tasks periodically synchronize and exchange data with each other
- *Barrier* = the point of synchronization
- MapReduce – a special case of BSP
 - Map Task = an operation applied to blocks of data in parallel
 - Reduce Task- when maps are “done” reduce the results to a single result
- MapReduce → open source: Hadoop



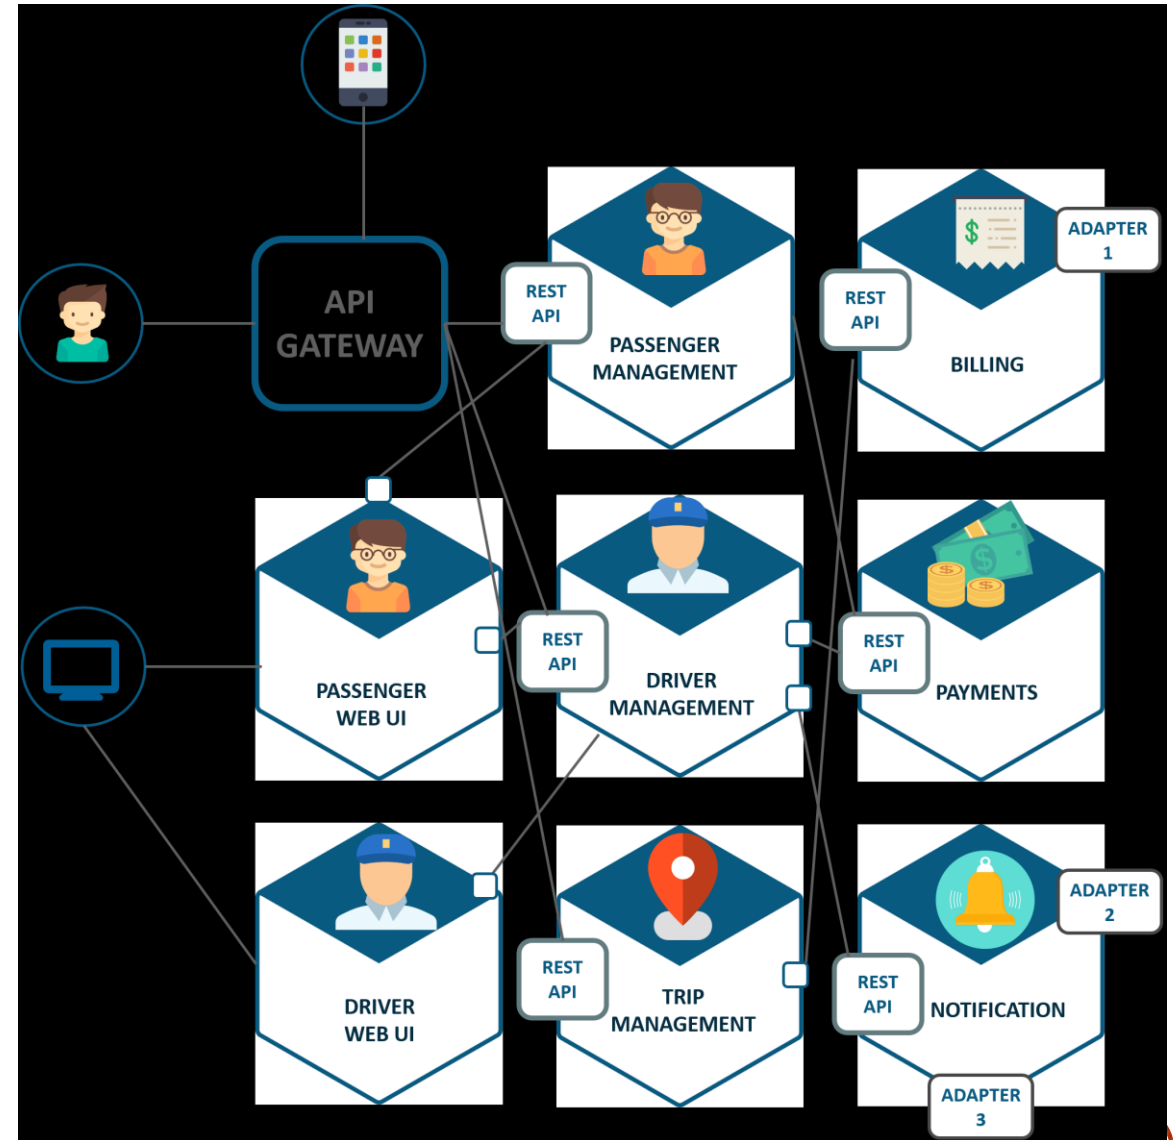
GRAPH PARALLEL

- Each function node represents a parallel invocation of the function on the distributed data structure
- The data is in distributed arrays or streams.
- build a data flow graph of the algorithms functions.
- The graph is compiled into parallel operators that are applied to the distributed data structures
- Data Analytic: Spark, Spark Streaming, Apache Flink, Storm, Google DataFlow
- Machine Learning: Google TensorFlow, MS Cognitive Toolkit



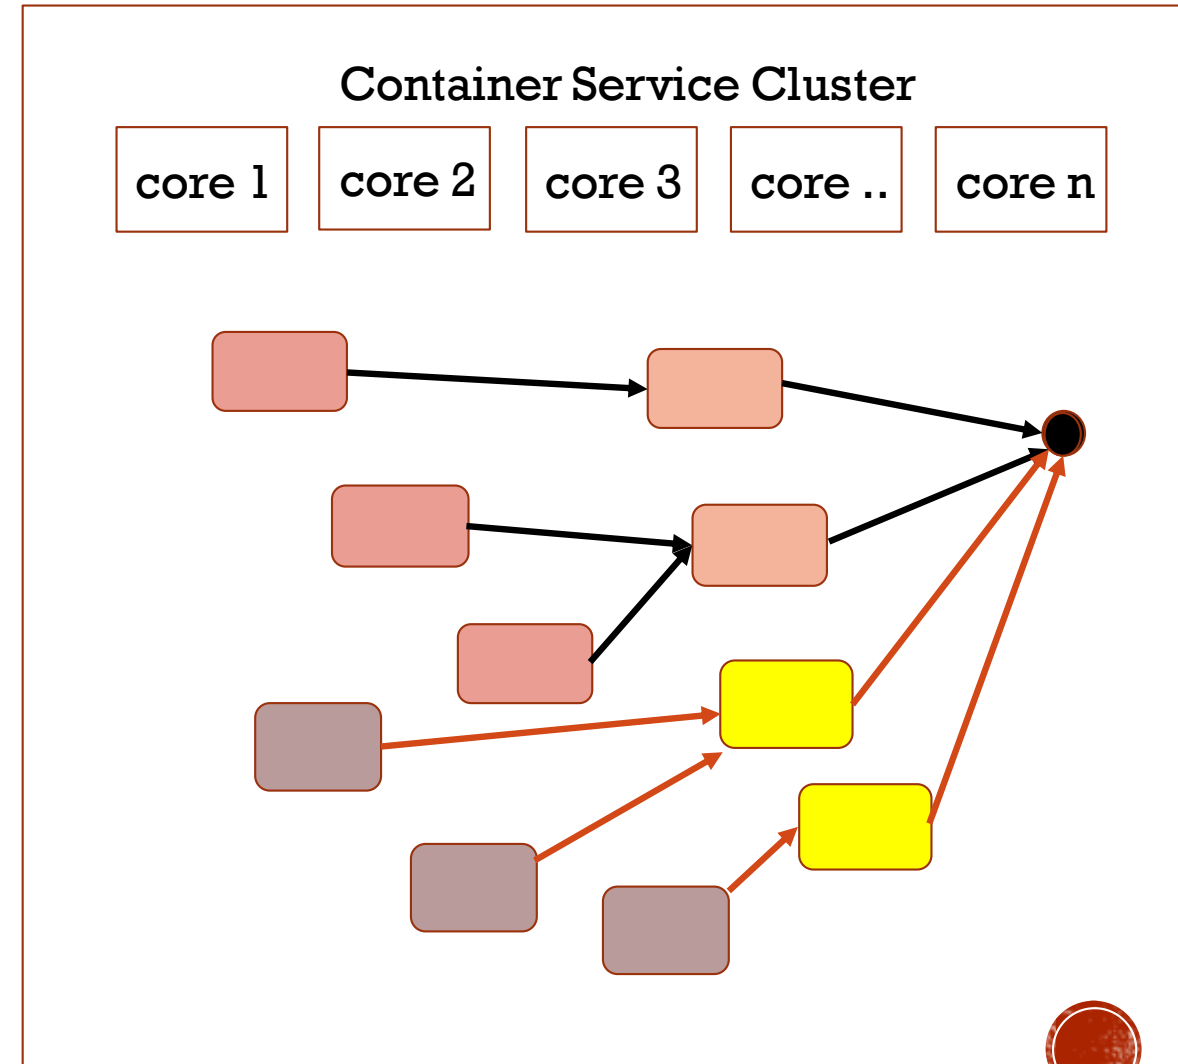
MICROSERVICES

- Divide a computation into small, mostly stateless components that can be
 - Easily replicated for scale
 - Communicate with simple protocols
 - Computation is as a swarm of communicating workers.



MICROSERVICES IN THE CLOUD

- Typically run as containers using a service deployment and management service
 - Amazon Elastic Container Service
 - Google Kubernetes
 - DCOS from Berkeley/Mesosphere
 - Docker Swarm




WHAT IS REST API?

- Interface for **distributed hypermedia systems**, created by Roy Fielding in 2000
- Guiding Principles of REST
 - Client – Server
 - Stateless
 - Cacheable
 - Uniform Interface
 - Layered System
 - Code on demand (optional)
- The key abstraction of information in REST is a **resource**
 - Identified by a resource identifier, ie URI
- Resources can be retrieved or transformed to another state by a set of methods
 - GET/PUT/POST/DELETE/PATCH
- The clients and servers exchange representations of resources by using a standardized interface and protocol typically HTTP



SERVERLESS

- New paradigm for service delivery
 - User provides a simple function to execute, against under certain conditions
- Very lightweight process, similar to daemon waiting for event to occur
- Managed by cloud infrastructure
 - No server management
 - Pay only while your code runs
 - Scale automatically
 - Highly available and fault tolerant
- Examples: AWS Lambda, Google Cloud Function, Azure Function



A SIMPLE SERVERLESS EXAMPLE - GCP

```
/**
 * HTTP Cloud Function.
 * This function is exported by index.js, and is executed when
 * you make an HTTP request to the deployed function's endpoint.
 *
 * @param {Object} req Cloud Function request context.
 *                  More info: https://expressjs.com/en/api.html#req
 * @param {Object} res Cloud Function response context.
 *                  More info: https://expressjs.com/en/api.html#res
 */
exports.helloGET = (req, res) => {
  res.send('Hello World!');
};
```

Have a simple function
that return “Hello World”

```
gcloud functions deploy helloGET --runtime nodejs6 --trigger-http
```

Deploy the function with
HTTP trigger

```
gcloud functions describe helloGET
```

```
https://GCP_REGION-PROJECT_ID.cloudfunctions.net/helloGET
```

Visit this URL in your browser.
You should see a Hello World!



SERVERLESS

- Pros
 - No need to write or manage backend code
 - Achieve Event-based programming without the complexity of building and maintaining the infrastructure.
- Cons
 - Vendor Lock-in
 - Architectural Complexity
 - AWS Lambda limits the number of concurrent executions you can be running of all your lambdas.
 - Startup Latency in FaaS
 - It takes time to initialize an instance of a function before each event.

