


HDFS — CONT'D



BASIC FILESYSTEM OPERATIONS

- **hadoop fs -help** – list all the filesystem command
- Make directory
hadoop fs -mkdir /user/name/testdir
- Copy a file from the local filesystem to HDFS
hadoop fs -copyFromLocal example.txt /user/name/.
hadoop fs -put example.txt /user/name/.
- Copy a file from HDFS to local
hadoop fs -copyToLocal /user/name/example.txt .
hadoop fs -get /user/name/example.txt .
- List files
hadoop fs -ls /user/name/

HDFS FILE LISTING

- Create a directory and list its content

```
% hadoop fs -mkdir books
```

```
% hadoop fs -ls .
```

Replication factor

Found 2 items

```
drwxr-xr-x - name supergroup 0 2009-04-02 22:41 /user/name/books  
-rw-r--r-- 1 name supergroup 18 2009-04-02 22:29 /user/name/example.txt
```

Replication factor determines how many times the file is replicated

- Empty for directories since the concept of replication does not apply
- Directories are treated as metadata and stored by the namenode, not the datanodes.



OTHER HADOOP SUPPORTED FILE SYSTEMS



BEYOND HDFS

HADOOP SUPPORTED FILE SYSTEMS

- Hadoop has an abstract notion of filesystem, of which HDFS is just one implementation

Filesystem	Java Implementation	Description
HDFS	<code>hdfs.DistributedFileSystem</code>	Hadoop's distributed filesystem. HDFS is designed to work efficiently in conjunction with MapReduce.
Local <i>file</i>	<code>fs.LocalFileSystem</code>	A filesystem for a locally connected disk with client-side checksums. Use <code>RawLocalFileSystem</code> for a local filesystem with no checksums.
WebHDFS	<code>hdfs.web.WebHdfsFileSystem</code>	A filesystem providing authenticated read/write access to HDFS over HTTP.

HADOOP SUPPORTED FILE SYSTEMS

Filesystem	Java Implementation	Description
FTP	<code>hdfs.ftp.FTPFileSystem</code>	A filesystem backed by a FTP server
S3	<code>fs.s3a.S3AFileSystem</code>	A filesystem backed by Amazon S3
Azure	<code>fs.azure.NativeAzureFileSystem</code>	A filesystem backed by Microsoft Azure
Swift	<code>fs.swift.snative.SwiftNativeFileSystem</code>	A filesystem backed by OpenStack Swift.


REFERENCES

- Hadoop The Definitive Guide 4th Edition,
Tom White.

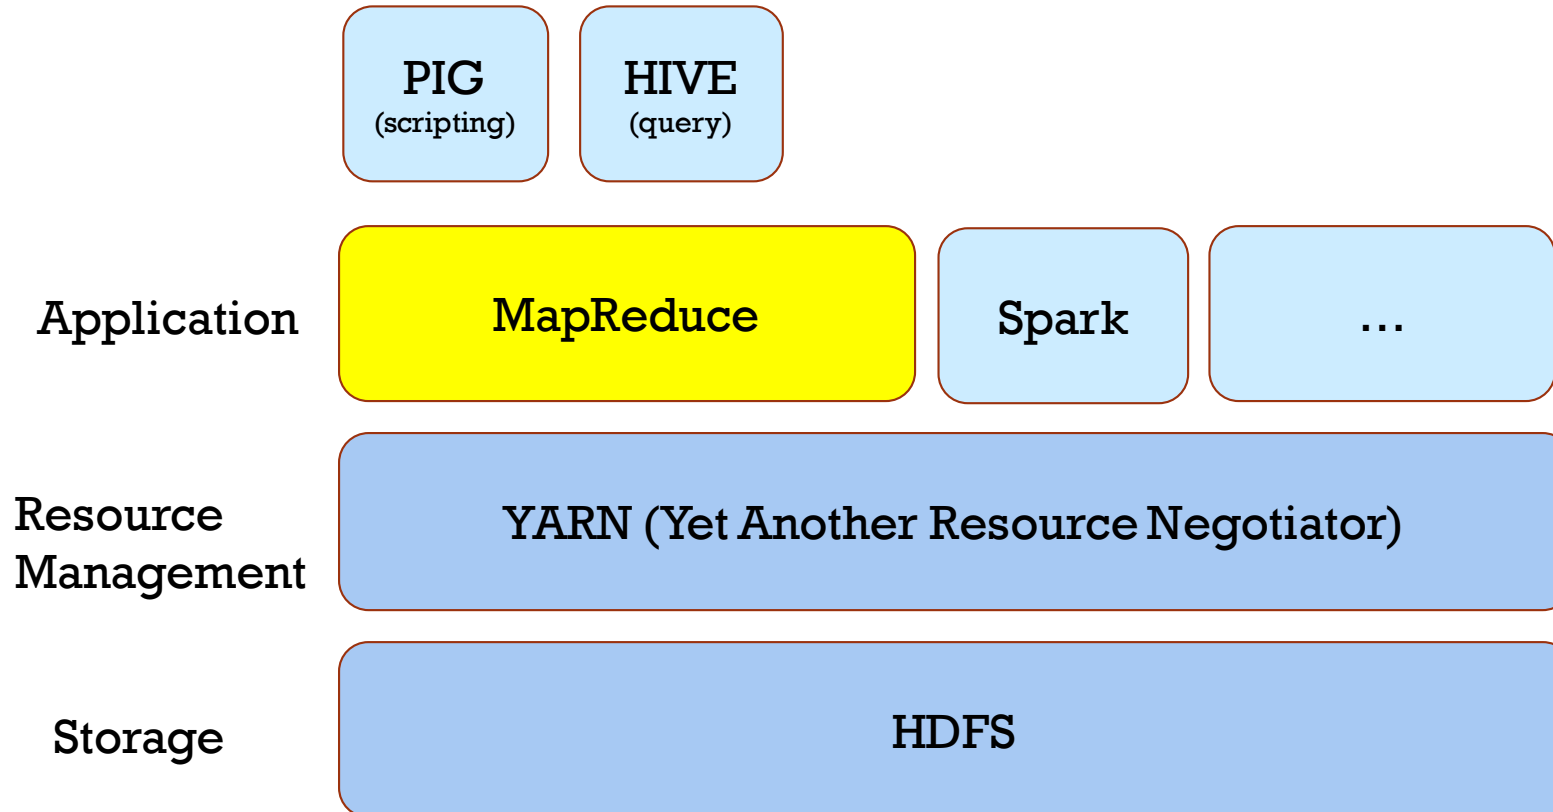
HDFS COMMON INTERVIEW QUESTIONS

- Explain the HDFS Architecture and list the various HDFS daemons in HDFS cluster?
- What is a NameNode in Hadoop?
- What is a DataNode?
- Is Namenode machine same as DataNode machine as in terms of hardware?
- What is the difference between traditional RDBMS and Hadoop?
- What is Secondary NameNode? Is it a substitute or back up node for the NameNode?
- What is checkpointing in Hadoop?
 - Combining Edit Logs and FsImage on the Secondary NameNode

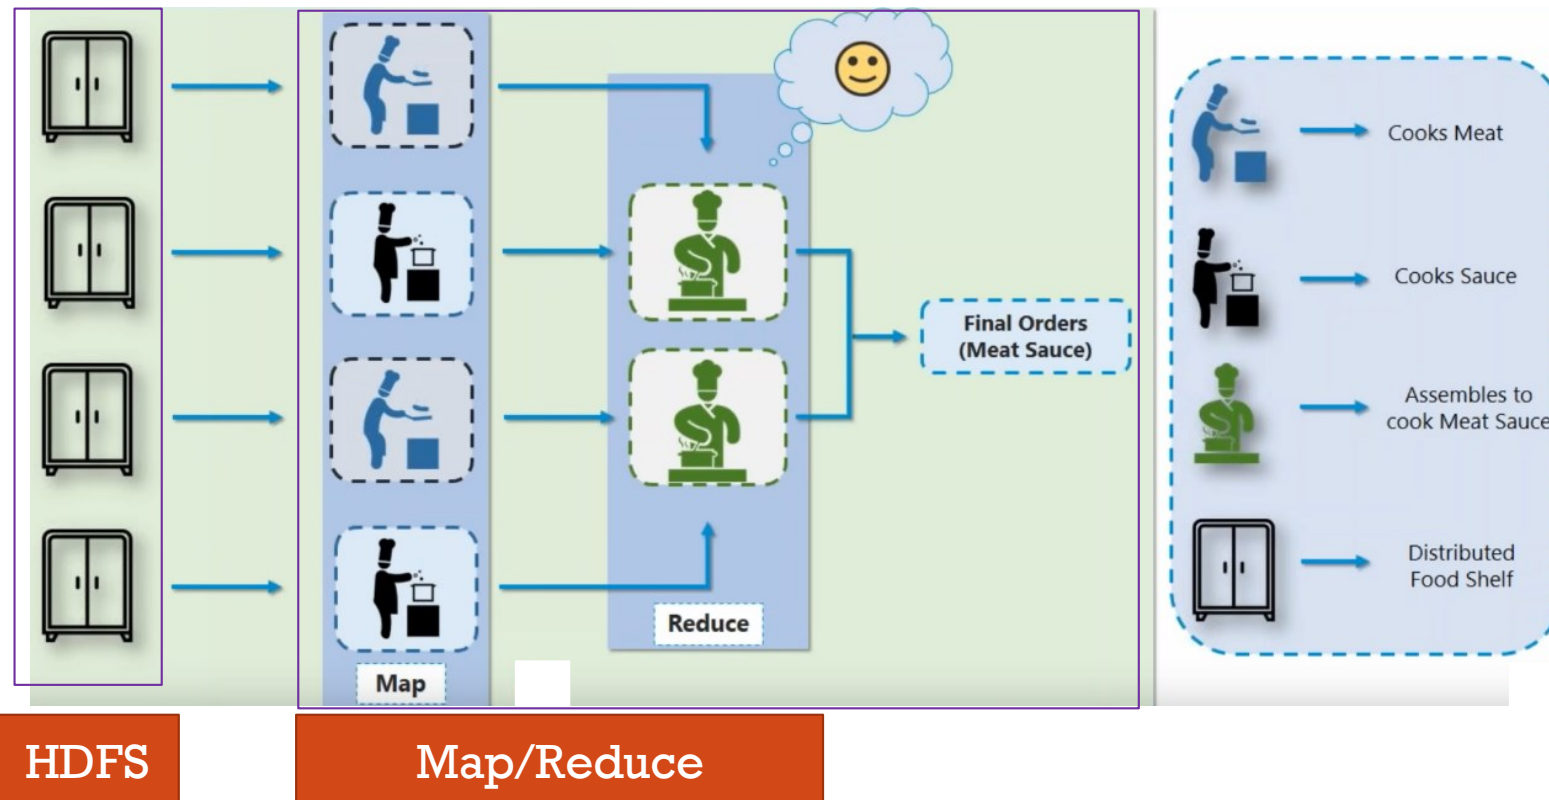
HADOOP MAPREDUCE



HADOOP SOFTWARE ARCHITECTURE



DO YOU REMEMBER?

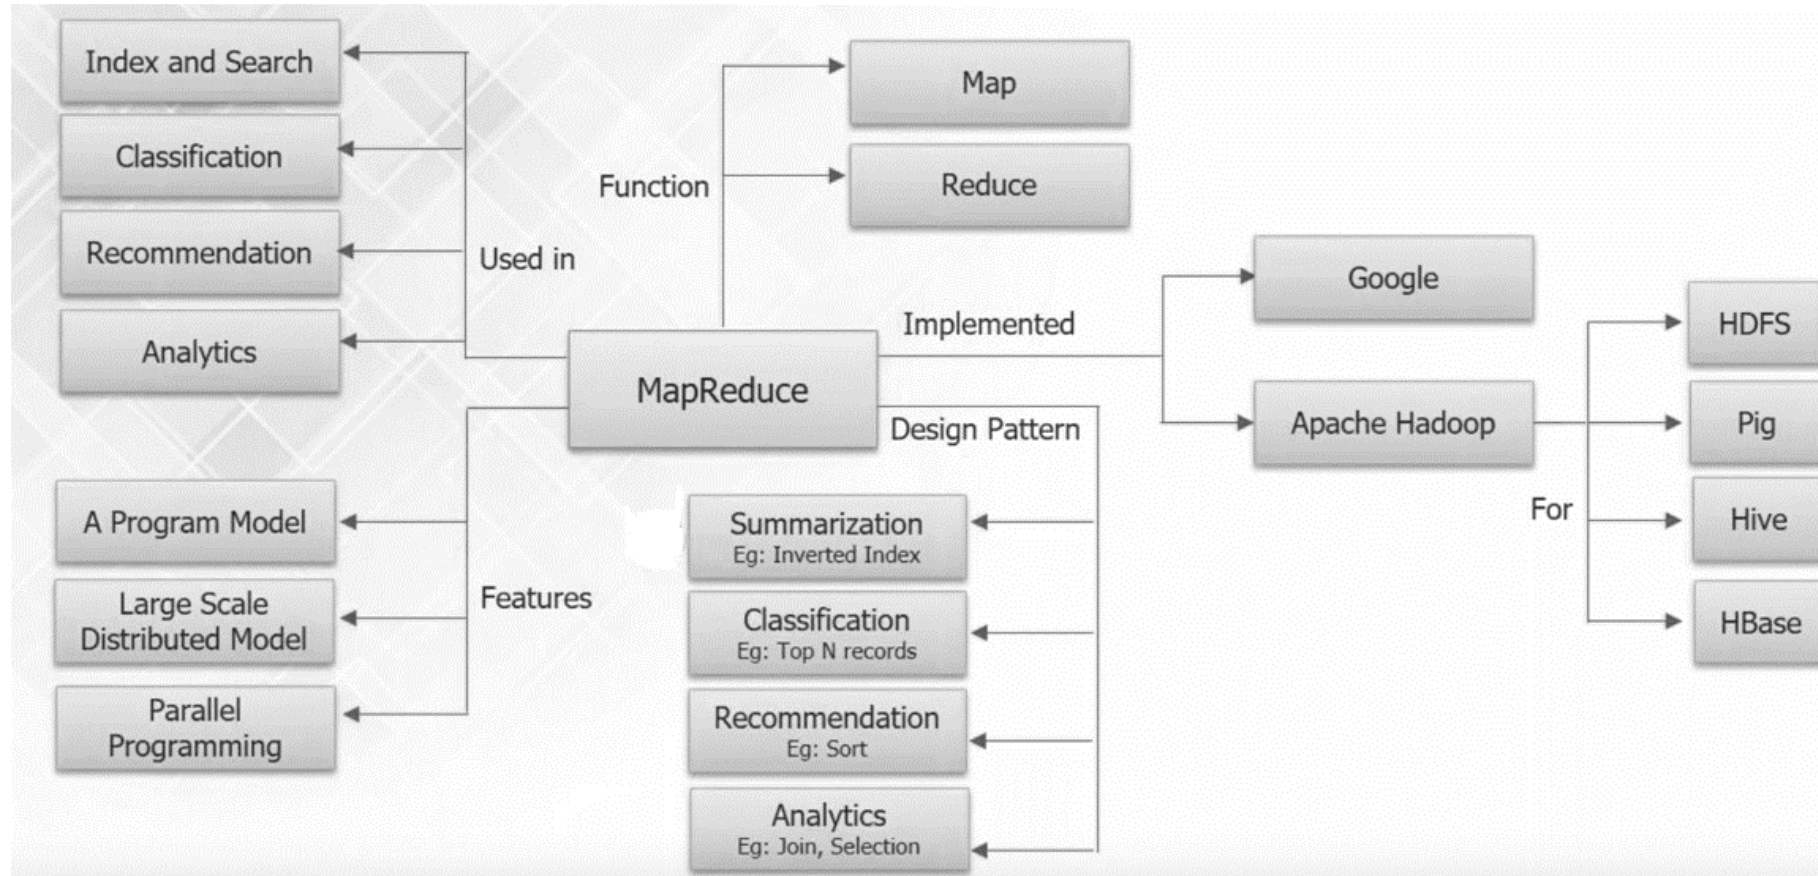


MAPREDUCE

- A programming model and an associated implementation for processing and generating large data sets with a parallel, distributed algorithm on a **Hadoop** cluster.
- It abstracts the problem from disk reads and writes, transforming it into a computation over sets for **keys and values**
- It works as a **batch query processor** that can run an ad hoc query against your whole data set and get result in a reasonable time



MAPREDUCE - SIMPLIFIED



WORD COUNT EXAMPLE

I can not do everything, but
still I can do something; and
because I cannot do
everything, I will not refuse
to do something I can do



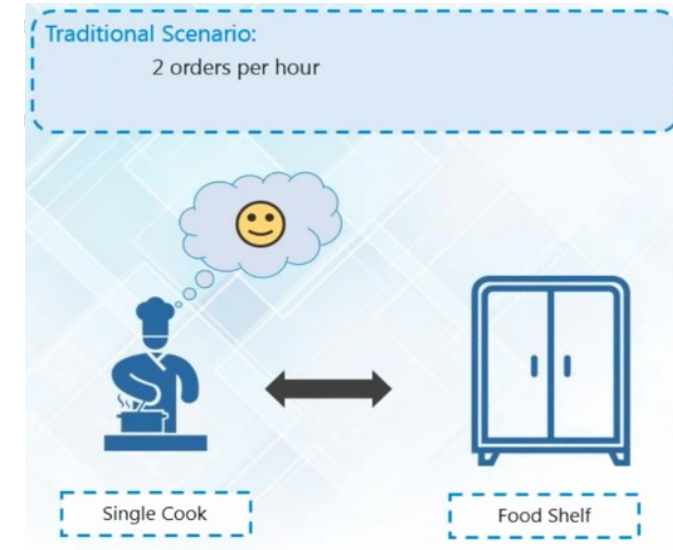
Word	Count
And	1
Because	1
But	1
Can	4
Do	4
Everything	2
I	5
Not	3
Refuse	1
Something	2
Still	1
To	1
Will	1



WORDCOUNT PROGRAM —

```
Define WordCount as Multiset;  
For Each Document in DocumentSet {  
    T = tokenize(document);  
    For Each Token in T {  
        WordCount[token]++;  
    }  
}  
Display(WordCount);
```

**Program Does NOT Scale for
Large Number of Documents**



WORDCOUNT PROGRAM – II

- A two-phased program
 - Distribute the work over several machines
 - Combine the outcome from each machine into the final word count
- Phase I – Document Processing
 - Each machine will process a fraction of the document set
- Phase II – Count Aggregation
 - Partial word counts from individual machines are combined into the final word count



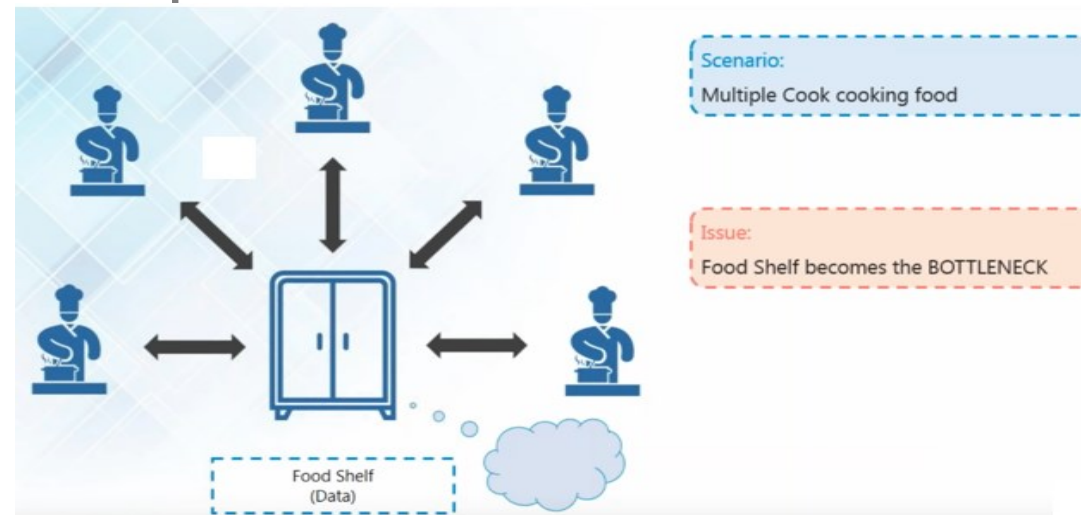
WORDCOUNT PROGRAM – II

Phase I

```
Define WordCount as Multiset;  
For Each Document in DocumentSubset {  
    T = tokenize(document);  
    For Each Token in T {  
        WordCount[token]++;  
    }  
}  
SendToSecondPhase(wordCount);
```

Phase II

```
Define TotalWordCount as Multiset;  
for each WordCount Received From  
firstPhase {  
    MultisetAdd (TotalWordCount,  
        WordCount);  
}
```



WORDCOUNT PROGRAM II — LIMITATIONS

- The program does not take into consideration the location of the documents
 - Storage server can become a bottleneck, if enough bandwidth is not available
- Storing WordCount and TotalWordCount in the memory is a flaw
 - When processing large document sets, the number of unique words can exceed the RAM capacity
- In Phase II, the aggregation machine becomes the bottleneck



WORDCOUNT PROGRAM

- HADOOP SOLUTION -

- Executing both phases in a distributed fashion on a cluster of machines that can run independently
- To achieve this, functionalities must be added
 - Store data over a cluster of processing machines
 - Partition intermediate data across multiple machines
 - Shuffle the partitions to the appropriate machines



GCP DEMO

- Redeem your coupon as shown in the Tutorials
- Install gsutil https://cloud.google.com/storage/docs/gsutil_install
- Navigate to your GCP Account at <https://console.cloud.google.com/>
- Create Your first cluster from Dataproc.
- Enjoy HDFS by connecting via SSH or Web Console to your master node

