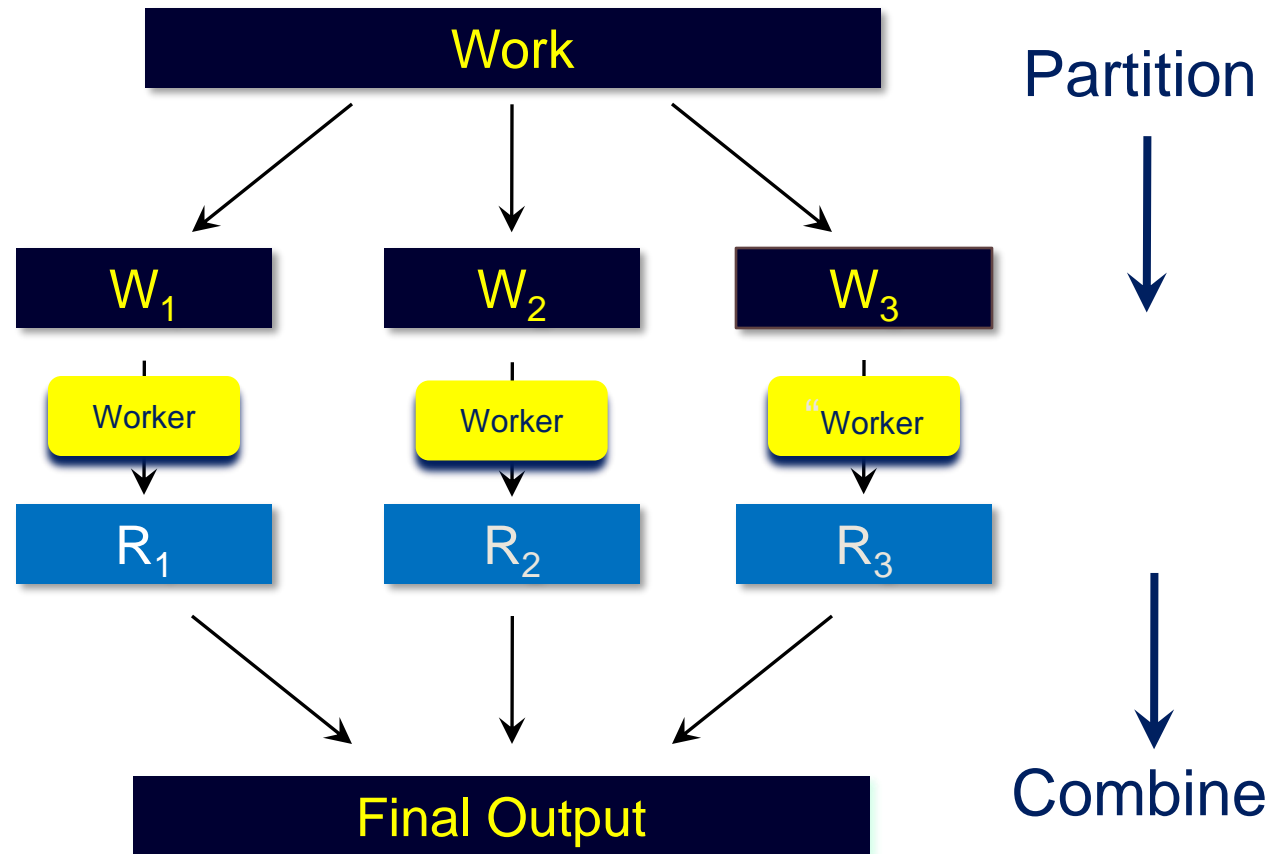


# MAPREDUCE — CONT'D



# DIVIDE AND CONQUER



# PARALLELIZATION CHALLENGES

- How do we assign work units to workers?
- What if we have more work units than workers?
- What if workers need to share partial results?
- How do we aggregate partial results?
- How do we know all the workers have finished?
- What if workers fail?



# DISTRIBUTED WORKERS COORDINATION

- Coordinating a large number of workers in a distributed environment is challenging
  - The order in which workers run may be unknown
  - The order in which workers interrupt each other may be unknown
  - The order in which workers access shared data may be unknown



# MAPREDUCE

## DATA-INTENSIVE PROGRAMMING MODEL

- Users specify the computation in terms of a **map()** and a **reduce()** function
  - **map**  $(k, v) \rightarrow \langle k', v' \rangle^*$
  - **reduce**  $(k', v') \rightarrow \langle k', v' \rangle^*$
- Underlying runtime system (YARN + HDFS)
  - Automatically parallelizes the computation across large-scale clusters of machines
  - Handles machine failures, communications and performance issues.



# WHAT ARE MAP AND REDUCE FUNCTIONS?

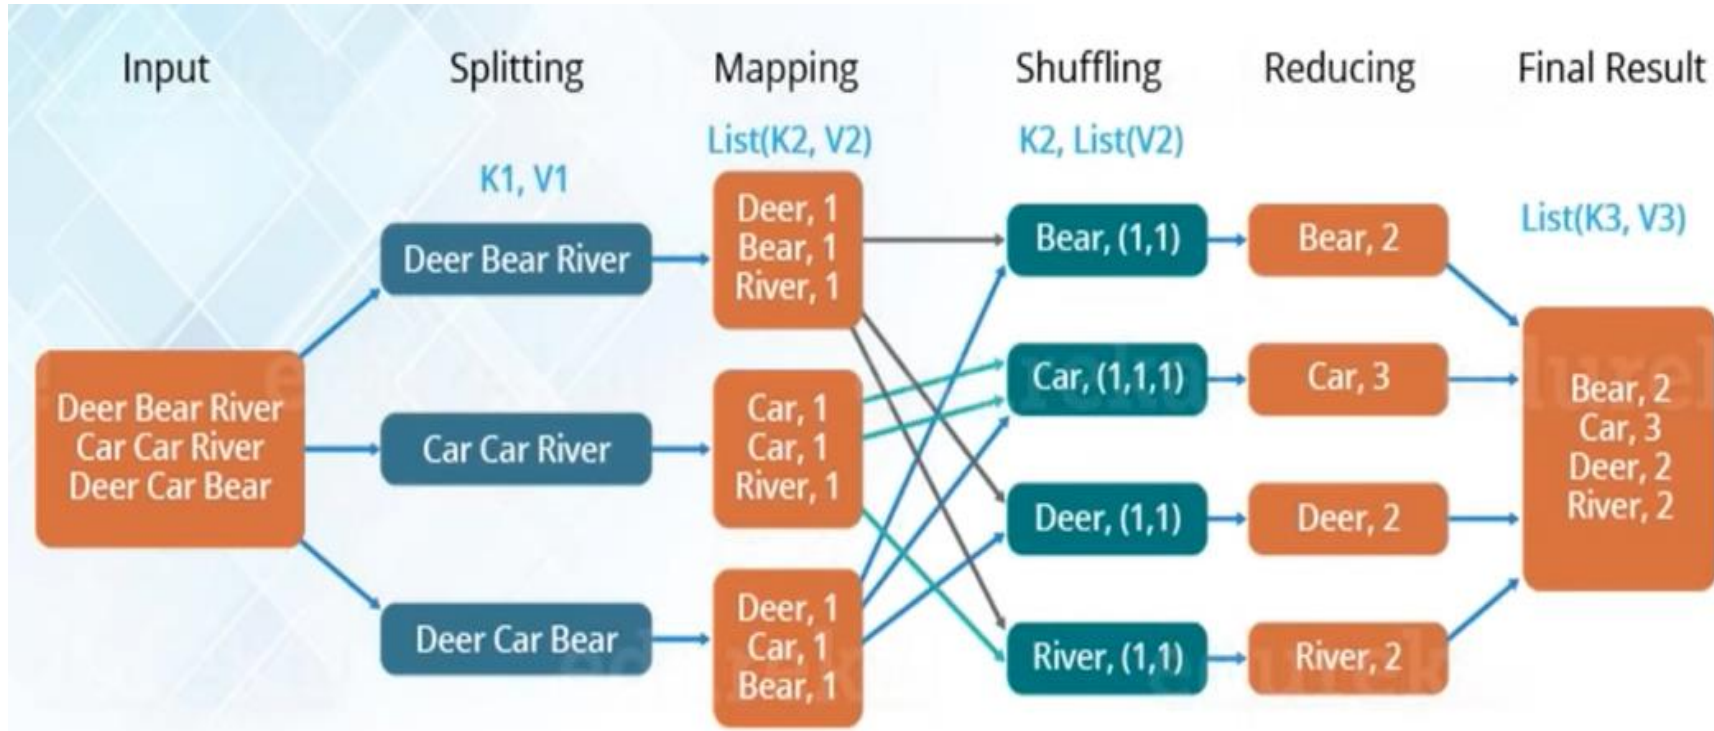
- At a high-level of abstraction, MapReduce codifies a generic “recipe” for processing large data set
    - Iterate over a large number of records
    - Extract something of interest from each
    - **Shuffle and sort intermediate results**
    - Aggregate intermediate results
    - Generate final output
- } Map
- } Shuffle
- } Reduce

Basic Tenet of MapReduce is Enabling a Functional Abstraction for the Map() and Reduce() operations

By default, Hadoop will take care of the Shuffle process for you!



# WORDCOUNT EXAMPLE IN HADOOP MAPREDUCE



# MAPREDUCE “RUNTIME” BASIC FUNCTIONS

(PROVIDED BY HADOOP)

- Handles scheduling
  - Assigns workers to map and reduce tasks
- Handles “process distribution”
  - Moves **processes to data**, not data to processes
- Handles synchronization among workers
  - Gathers, sorts, and shuffles intermediate data
- Handles errors and faults, dynamically
  - Detects worker failures and restarts





# MAPREDUCE DATA FLOW

- A MapReduce job is a unit of work to be performed
  - Job consists of the **MapReduce Program**, the **Input data** and the **Configuration Information**
- The MapReduce job is divided it into two types of tasks – *map* tasks and *reduce* tasks
- The Input data is divided into fixed-size pieces called *splits*
  - One map task is created for each split
  - The map function is run on each split
- Configuration information indicates where the input lies and the output is stored

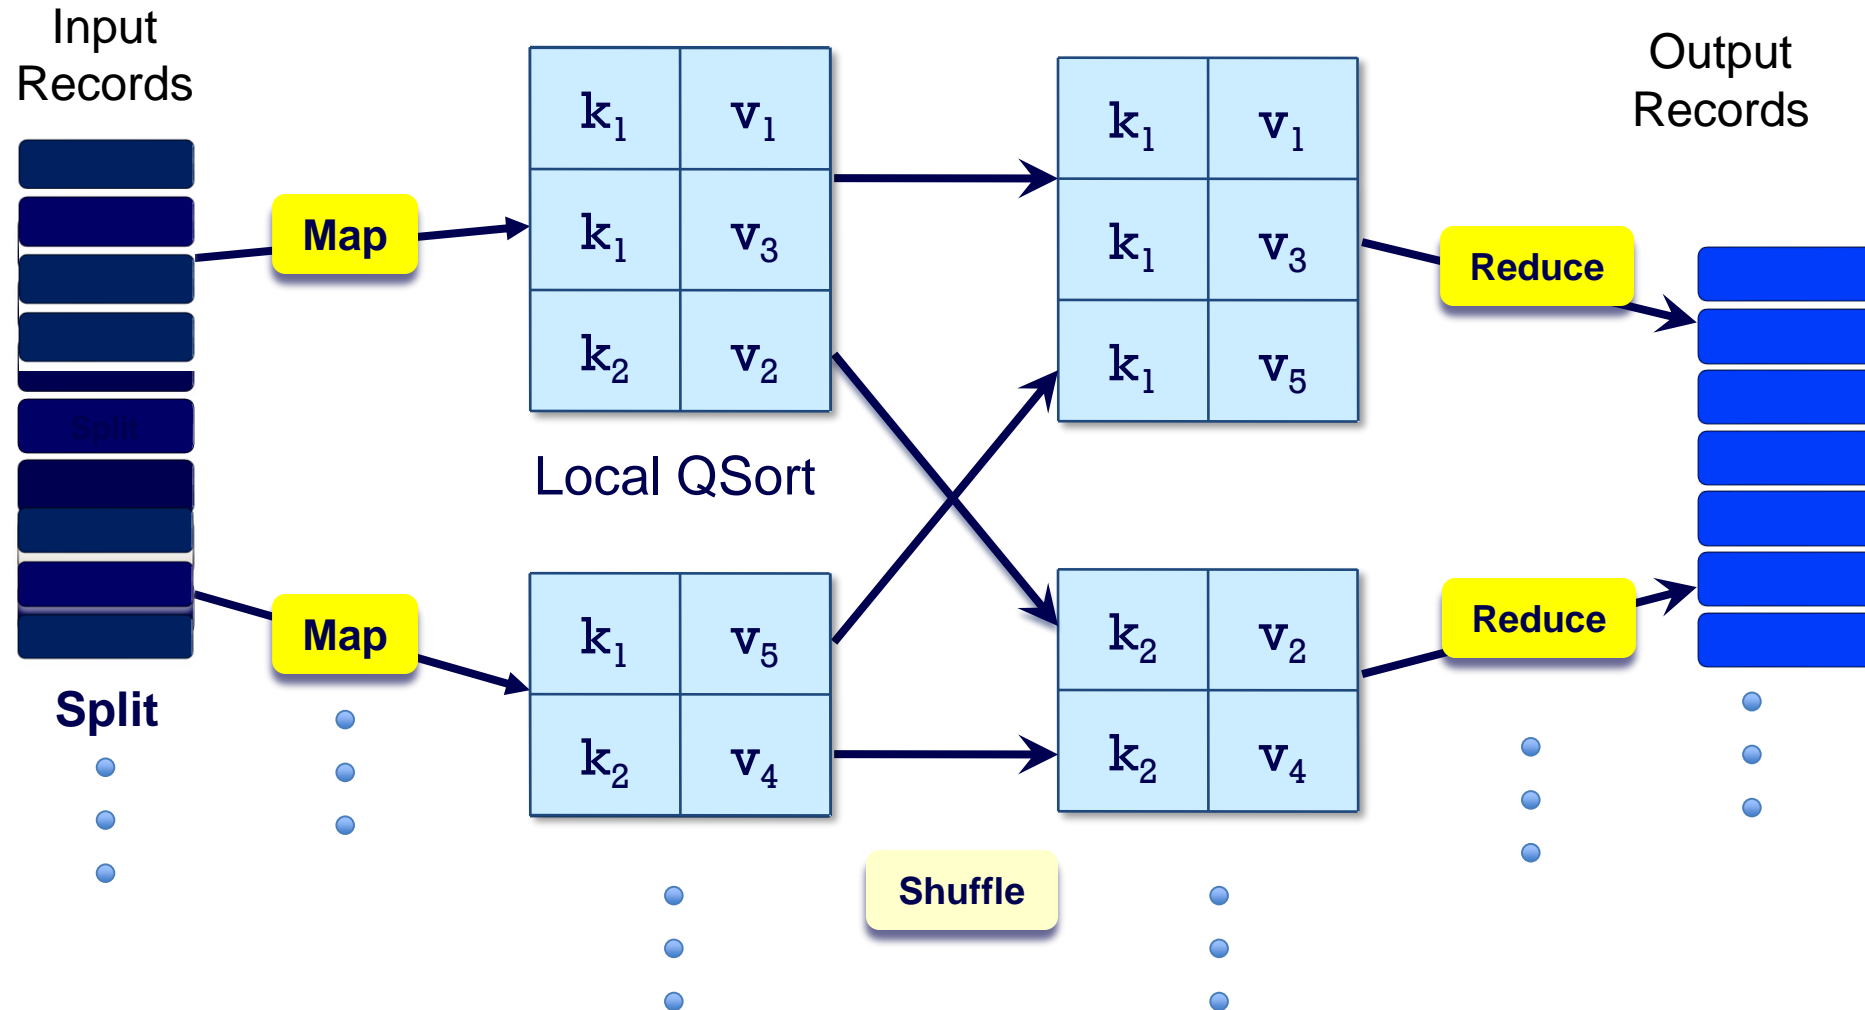


# FROM MAP TO REDUCE

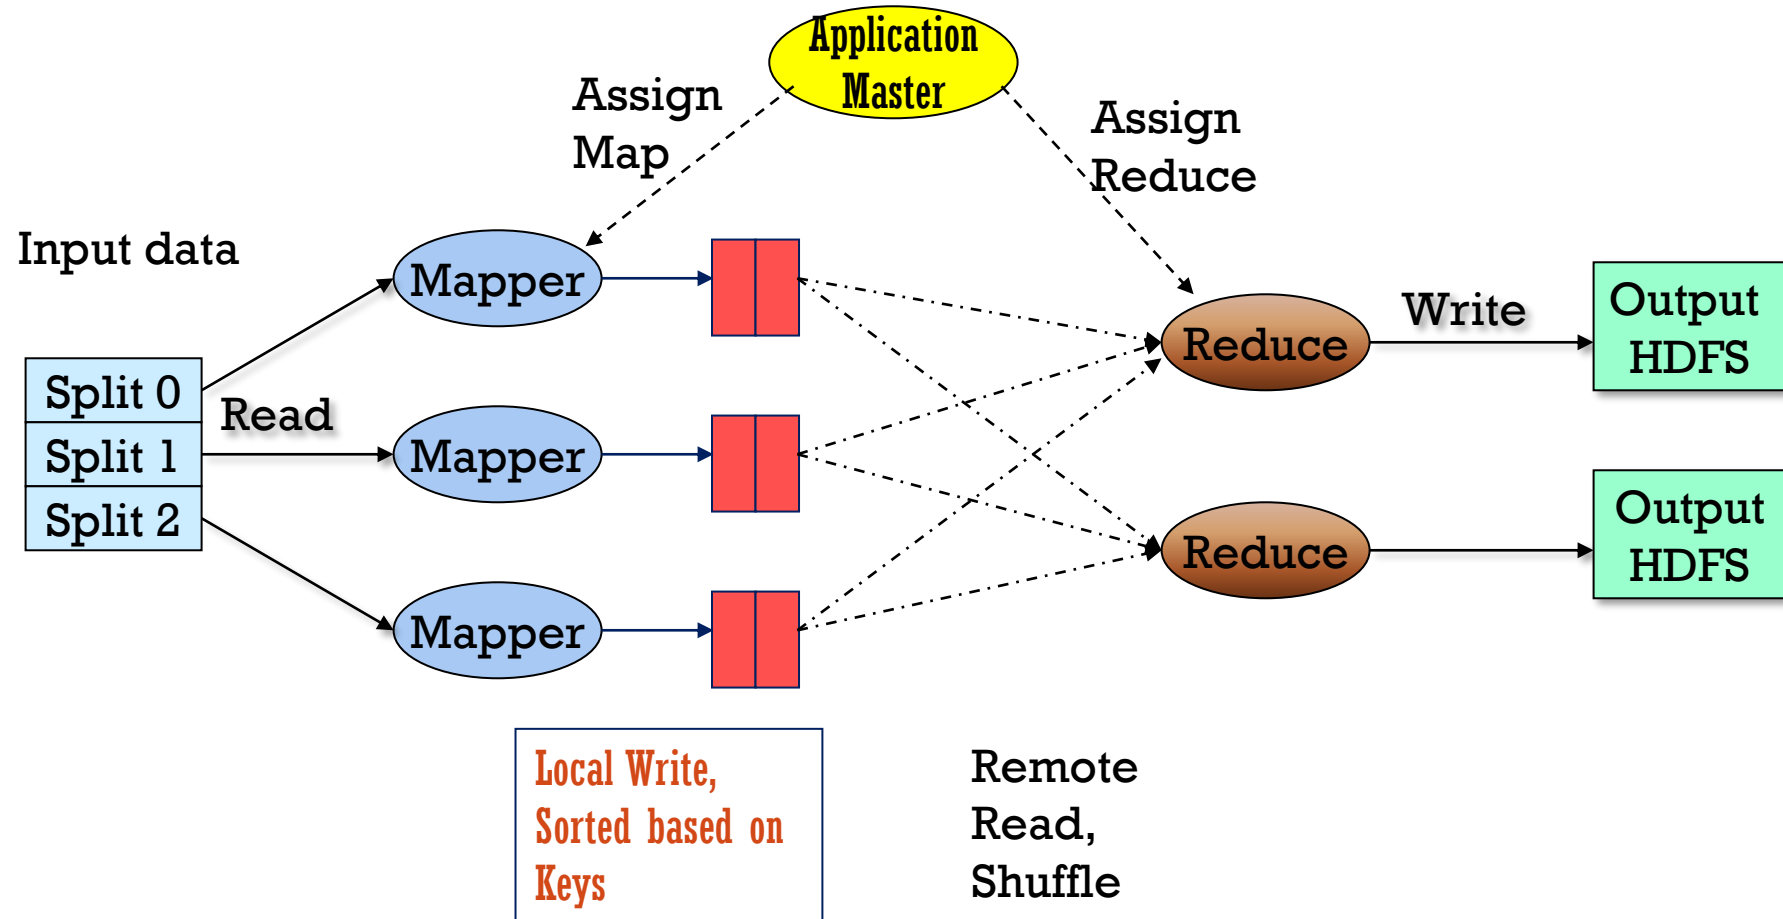
- Select a number  $R$  of reduce tasks.
- Divide the intermediate keys into  $R$  groups,
  - Use an efficient hashing function
- Each Map task creates, at its own processor,  $R$  files of intermediate key-value pairs, sorted by key, and one for each Reduce task
- Its output file is transfer to the Reducer's container
- Once all Mapper tasks results are transferred to the Reducer, it merges the results and compute the final output
- The Reducer then writes the output to HDFS



# MAPREDUCE DATA FLOW



# DISTRIBUTED EXECUTION OVERVIEW



# READING

- <https://www.guru99.com/introduction-to-mapreduce.html>

